



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# Porting Applications to GPUs

**CSCS Summer School 2025**

**Andreas Jocksch, Prashanth Kanduri, Radim Janalik & Ben Cumming**

# Using GPUs in your Application?

Rule #1: **DO NOT** develop your own GPU code

# Libraries

There are numerous libraries available for GPUs. Some from Nvidia and many open source.

- [cuBLAS](#): Dense linear algebra primitives
- [Thrust](#): C++ STL-like algorithms and containers
- [cuRAND](#) and [Random123](#): Random numbers
- [cuFFT](#): Fast-Fourier Transforms
- [Kokkos](#): Generic performance portable parallel motifs

... and many more!

Take some time to investigate what is available before starting.

# Are you going to write your own code?

## Directives

OpenACC and OpenMP define *directives* that can be used to instruct the compiler how to generate GPU code.

In theory the easiest path for porting.

## GPU-Specific Languages

Languages designed for GPU programming.

Maximum flexibility and performance.

For example: CUDA, OpenCL and SYCL.

# Things to Consider

Before starting on a GPU implementation, it pays to ask some questions and do some preliminary exploration:

1. Is my program computationally or bandwidth intensive?
  2. Does it have enough parallel work to utilize the GPU?
  3. Must I change algorithms to expose enough parallelism?
  4. Are there serial bottlenecks that will limit scaling? (see Amdahl's and Gustafson's laws)
  5. Is the pain worth the gain?
- Questions 1, 2 and 3 will be discussed in this course.
  - Questions 5 requires answers for 1-4.

# CUDA

CUDA (Compute Unified Device Architecture) is a **parallel computing platform** and **API** for CUDA-enabled Nvidia GPUs

We use CUDA as short hand for CUDA C/C++ and API

- CUDA C++ is a **superset** of C++
- Adds keywords for writing kernels to run on the GPU
- Adds syntax for launching kernels on the GPU

The CUDA toolkit is more than a programming language. It also provides:

- Runtime API for managing GPU resources and execution
- Tools including profilers and debuggers

# Compiling CUDA

CUDA code is compiled with the `nvcc` compiler driver

- source files have `.cu` extension
- headers have `.h`, `.hpp`, `.hcu` extension

CUDA compilation involves multiple splitting, compilation, preprocessing and merging steps

- `nvcc` hides this complexity from the user
- It closely mimics the interface of the GNU compiler
- Behind the scenes it:
  - uses GCC to compile the code that runs on CPU;
  - and compiles the GPU code separately

# Compiling CUDA

Example CUDA compilation

```
> nvcc -arch=sm_90 -lineinfo -O2 -std=c++11 -g -o foo foo.cu
```

Some flags are for **device** code generation:

- **-arch=sm\_90** target GPU architecture (Hopper streaming multiprocessor)
- **-lineinfo** debug information for device code

Some are for **host**:

- **-g** debug information for host code

And some are for both **host** and **device**:

- **-O2** optimization level
- **-std=c++11** target language
- **-o foo** name of executable

# Exercise: Getting Started on Daint (Alps)

In this exercise we will get introduced to Daint and make sure that everybody is set up.

```
# log on to daint (from ela.csccs.ch) with your course username & password
> ssh -X <your account name>@daint

# go to scratch and get the course material
> cd $SCRATCH/SummerUniversity
> git pull

# compile and test the demo
> cd cuda/practicals/demos
> cat hello.cu
> uenv image pull prgenv-gnu/24.11:v2
> uenv start --view=default prgenv-gnu/24.11:v2
> nvcc -arch=sm_90 hello.cu -o hello
# run the job on the course reservation for maximum 60 minutes
# find your group id with the "groups" command
> srun -N 1 -A class05 --reservation=summer_uni -t60 ./hello
```